

Assignment 2 - Stock Data Analysis

Barry Li, Bella Batkin, Emma Torossian, Maxim Glietz, Robin Dreher

2025-10-26

Table of contents

| | |
|---|-----------|
| Part 1 Questions & Answers | 3 |
| 1. How many unique tickers are in your data? | 3 |
| 2. How many unique companies are in your data? | 4 |
| 3. Display the top 5 companies by largest mean trading volume, in a table. | 4 |
| 4. Display the total trading volume of the top 3 exchanges (table). | 5 |
| 5. Visualise the total trading volume of the top 3 exchanges (bar plot). | 5 |
| 6. How many companies have more than one ticker? | 6 |
| 7. Which ticker has the largest positive mean return (simple daily return)? | 7 |
| 8. Which company has the largest positive mean return (simple daily return)? | 8 |
| 9. Which industry is represented by the most companies? | 8 |
| Part 2 Extended Analysis | 9 |
| 1. Calculate simple weekly returns for each ticker in the full dataset | 9 |
| 2. Categorise your data into decile groups | 9 |
| 3. Display a table showing the top ticker in each decile group | 10 |
| 4. Select the top ticker from the 60% decile group | 11 |
| 5. Plot the autocorrelation function for this ticker's entire set of weekly returns | 12 |
| Part 3 Regression - Fama-French 3 Factor Model | 12 |
| 1. Load and clean the weekly Fama-French 3 factor data | 12 |
| 2. Fit the Fama-French 3 factor model to the weekly returns of the stock in Part 2 | 13 |
| Analysis Summary | 14 |

List of Figures

| | |
|---|----|
| 1 Total Trading Volume of the Top 3 Exchanges | 6 |
| 2 Autocorrelation Function | 12 |

List of Tables

| | |
|--|----|
| 1 Compressed dataset | 3 |
| 4 Top 5 companies by largest mean trading volume | 4 |
| 5 Total trading volume of the top 3 exchanges | 5 |
| 6 Companies with more than one ticker | 6 |
| 10 Simple weekly returns for each ticker | 9 |
| 11 Data categorized by decile groups | 10 |
| 12 Top tickers by decile group | 11 |
| 14 Fama-French 3 Factor Model | 13 |

Before messing around with the stock data, the environment should install and load the dplyr and lubridate packages as well as others to perform easier data analysis. Additionally, we disable any warning messages for cleaner output. We also remove any rows with NA values in the prcod column.

```
library(dplyr)
library(readr)
library(data.table)
library(lubridate)
library(ggplot2)
options(warn=-1)

data = fread("compressed_data.csv.gz") %>%
  filter(!is.na(prcod)) %>%
  mutate(datadate = as.Date(datadate, "%d/%m/%Y"))
head(data |> select(-conm, -gvkey))
```

A data.table: 6 × 9

Table 1: Compressed dataset

| tic <chr> | datadate <date> | exchg <int> | sic <int> | cshtrd <dbl> | prccd <dbl> | prchd <dbl> | preld <dbl> | prcod <dbl> |
|--------------|--------------------|----------------|--------------|-----------------|----------------|----------------|----------------|----------------|
| PNW | 2023-01-03 | 11 | 4911 | 1442534 | 74.63 | 76.4125 | 73.380 | 76.25 |
| PNW | 2023-01-04 | 11 | 4911 | 954218 | 75.39 | 76.0950 | 74.630 | 75.10 |
| PNW | 2023-01-05 | 11 | 4911 | 994775 | 73.65 | 75.0950 | 73.305 | 74.88 |
| PNW | 2023-01-06 | 11 | 4911 | 729808 | 75.46 | 76.0200 | 74.480 | 74.49 |
| PNW | 2023-01-09 | 11 | 4911 | 656127 | 75.55 | 76.4800 | 75.240 | 75.24 |
| PNW | 2023-01-10 | 11 | 4911 | 763254 | 75.65 | 75.6950 | 74.880 | 75.31 |

Part 1 Questions & Answers

1. How many unique tickers are in your data?

```
summary_table_1 <- tibble(
  X = "Unique tickers",
  Y = length(unique(data$tic))
)
summary_table_1
```

```
cat("1. There are", length(unique(data$tic)), "unique tickers.")
```

A tibble: 1 × 2

| X <chr> | Y <int> |
|----------------|---------|
| Unique tickers | 502 |

1. There are 502 unique tickers.

2. How many unique companies are in your data?

```
summary_table_2 <- tibble(
  X = "Unique company names",
  Y = length(unique(data$conm))
)
summary_table_2
cat("\n2. There are", length(unique(data$conm)), "unique company names.")
```

A tibble: 1 × 2

| X <chr> | Y <int> |
|----------------------|---------|
| Unique company names | 499 |

2. There are 499 unique company names.

3. Display the top 5 companies by largest mean trading volume, in a table.

```
data_3 = data %>%
  group_by(tic) %>%
  summarise(mean_trading_v = mean(cshtrd, na.rm = TRUE)) %>%
  ungroup() %>%
  arrange(desc(mean_trading_v))
data_3[1:5,]
```

A tibble: 5 × 2

Table 4: Top 5 companies by largest mean trading volume

| tic <chr> | mean_trading_v <dbl> |
|-----------|----------------------|
| TSLA | 115314383 |
| NVDA | 113131835 |
| PLTR | 60056251 |
| AAPL | 57736403 |

| tic <chr> | mean_trading_v <dbl> |
|-----------|----------------------|
| AMD | 57143415 |

4. Display the total trading volume of the top 3 exchanges (table).

```
data_4 = data %>%
  group_by(exchg) %>%
  summarise(total_trading_v = sum(cshtrd, na.rm = TRUE)) %>%
  ungroup() %>%
  arrange(desc(total_trading_v))
data_4[1:3,]
```

A tibble: 3 × 2

Table 5: Total trading volume of the top 3 exchanges

| exchg <int> | total_trading_v <dbl> |
|-------------|-----------------------|
| 11 | 681415756062 |
| 14 | 570830885382 |
| 21 | 385399362 |

5. Visualise the total trading volume of the top 3 exchanges (bar plot).

```
ggplot(data_4, aes(x = as.character(exchg), y = total_trading_v/1000000)) +
  geom_bar(stat = "identity", color = "darkblue", fill = "darkblue") +
  geom_text(aes(label = round(total_trading_v/1000000)),
            vjust = -0.3,                      # position above the bar
            size = 5) +                         # text size
  labs(title = "Total Trading Volume of the Top 3 Exchanges",
       x= "exchange", y="Total Trading Volume in Millions") +
  theme(plot.title = element_text(hjust = 0.5))
```

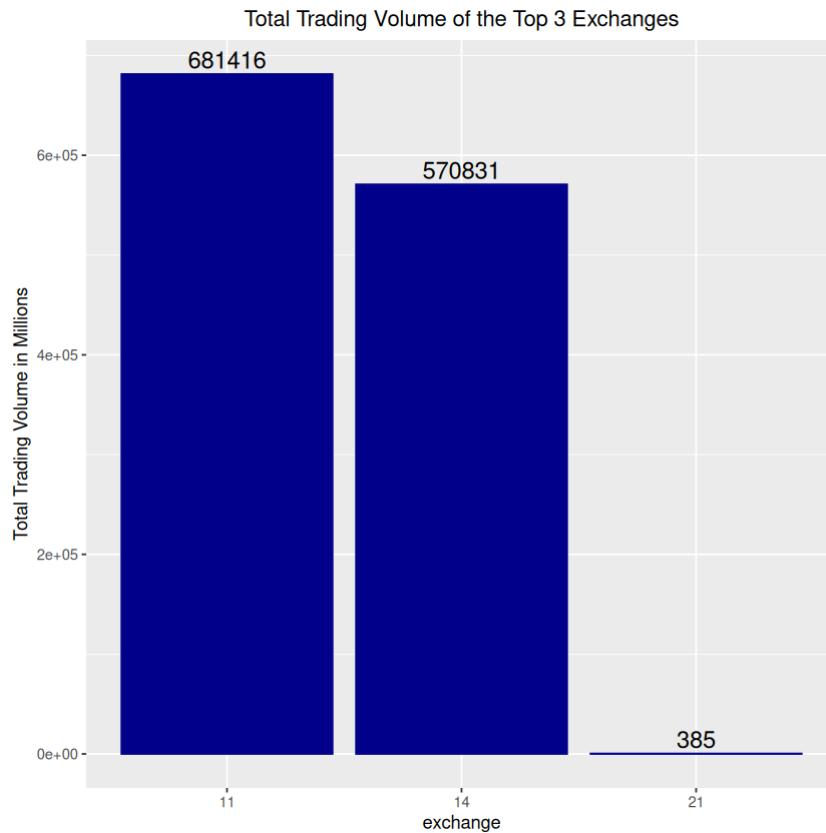


Figure 1: Total Trading Volume of the Top 3 Exchanges

6. How many companies have more than one ticker?

```
data_6 = data %>%
  group_by(conm) %>%
  summarise(no_of_tickers = n_distinct(tic)) %>%
  ungroup() %>%
  filter(!no_of_tickers == 1)
data_6[1:4,]
nr_companies = nrow(data_6)
cat("6. There are", nr_companies, "companies with more than one ticker.")
```

A tibble: 4 × 2

Table 6: Companies with more than one ticker

| conm <chr> | no_of_tickers <int> |
|--------------|---------------------|
| ALPHABET INC | 2 |
| FOX CORP | 2 |
| NEWS CORP | 2 |
| NA | NA |

6. There are 3 companies with more than one ticker.

7. Which ticker has the largest positive mean return (simple daily return)?

```
# 7. Which ticker has the largest positive mean return (simple daily return)?
data = data %>%
  group_by(tic) %>%
  mutate(return = prccd/lag(prccd)-1) %>%
  ungroup()

data_7 = data %>%
  group_by(tic) %>%
  summarise(mean_return = mean(return, na.rm = TRUE)) %>%
  ungroup() %>%
  arrange(desc(mean_return))

highest_mean_return = max(data_7$mean_return)

highest_mean_return_ticker = data_7$tic[
  which.max(data_7$mean_return)
]

summary_table_7 <- tibble(
  Ticker = highest_mean_return_ticker,
  Mean_return_perc = round(highest_mean_return, 4)*100
)
summary_table_7

cat("7. The", highest_mean_return_ticker,
  "ticker had the highest mean daily return.")
cat("\n-> The return was", round(highest_mean_return, 4)*100, "%.")
```

A tibble: 1 × 2

| Ticker <chr> | Mean_return_perc <dbl> |
|--------------|------------------------|
| PLTR | 0.58 |

7. The PLTR ticker had the highest mean daily return.

-> The return was 0.58 %.

8. Which company has the largest positive mean return (simple daily return)?

```
summary_table_8 <- tibble(
  Ticker = highest_mean_return_ticker,
  Company = highest_mean_return_company,
  Mean_return_perc = round(highest_mean_return, 4)*100
)
summary_table_8

highest_mean_return_company = data$comm[
  which(data$tic == highest_mean_return_ticker)][1]
]
cat("8. The", highest_mean_return_company,
"company had the highest mean daily return.")
```

A tibble: 1 × 3

| Ticker <chr> | Company <chr> | Mean_return_perc <dbl> |
|--------------|---------------------------|------------------------|
| PLTR | PALANTIR TECHNOLOG INC | 0.58 |

8. The PALANTIR TECHNOLOG INC company had the highest mean daily return.

9. Which industry is represented by the most companies?

```
data_9 = data %>%
  group_by(sic) %>%
  summarise(no_companies = n_distinct(comm)) %>%
  ungroup() %>%
  arrange(desc(no_companies))
most_represented_industry = data_9$sic[
  which.max(data_9$no_companies)
]
no_companies_in_most_represented_industry = max(data_9$no_companies)

summary_table_9 <- tibble(
  SIC = most_represented_industry,
  No_of_companies = no_companies_in_most_represented_industry
)
summary_table_9

cat("9. The", most_represented_industry,
"SIC industry has the most companies.")
cat("\n-> There are", no_companies_in_most_represented_industry,
"companies in that industry.")
```

A tibble: 1 × 2

| SIC <int> | No_of_companies <int> |
|-----------|-----------------------|
| 6798 | 28 |

9. The 6798 SIC industry has the most companies.
-> There are 28 companies in that industry.

Part 2 Extended Analysis

After preparing the data we carry out the following analysis.

1. Calculate simple weekly returns for each ticker in the full dataset

```
data_weekly = data %>%
  ## determine weekly returns based on fridays
  group_by(tic, datadate = floor_date(datadate, "week") + 5) %>%
  summarise(weekly_close = last(prccd)) %>%
  arrange(tic, datadate) %>%
  group_by(tic) %>%
  mutate(simple_w_r = (weekly_close / lag(weekly_close)) - 1) %>%
  ungroup()
head(data_weekly)
```

`summarise()` has grouped output by 'tic'. You can override using the `groups` argument.

A tibble: 6 × 4

Table 10: Simple weekly returns for each ticker

| tic <chr> | datadate <date> | weekly_close <dbl> | simple_w_r <dbl> |
|-----------|-----------------|--------------------|------------------|
| A | 2023-01-06 | 147.67 | NA |
| A | 2023-01-13 | 156.92 | 0.062639670 |
| A | 2023-01-20 | 155.92 | -0.006372674 |
| A | 2023-01-27 | 155.69 | -0.001475115 |
| A | 2023-02-03 | 154.55 | -0.007322243 |
| A | 2023-02-10 | 152.55 | -0.012940796 |

2. Categorise your data into decile groups

(We do not remove zero returns from the data).

```

c_breaks = seq(0.1, 1, by = 0.1)
c_labels = paste0((1:(length(c_breaks) - 1)) * 10, "%")

data_weekly_deciles <- data_weekly %>%
  filter(!is.na(simple_w_r)) %>%
  mutate(
    deciles = cut(
      simple_w_r,
      breaks = quantile(
        simple_w_r,
        probs = c_breaks,
        type = 9,
        na.rm = TRUE
      ),
      labels = c_labels,
      include.lowest = TRUE
    )
  ) %>%
  arrange(tic)
head(data_weekly_deciles)

```

A tibble: 6 × 5

Table 11: Data categorized by decile groups

| tic <chr> | date <date> | weekly_close <dbl> | simple_w_r <dbl> | deciles <fct> |
|-----------|-------------|--------------------|------------------|---------------|
| A | 2023-01-13 | 156.92 | 0.062639670 | 90% |
| A | 2023-01-20 | 155.92 | -0.006372674 | 30% |
| A | 2023-01-27 | 155.69 | -0.001475115 | 40% |
| A | 2023-02-03 | 154.55 | -0.007322243 | 30% |
| A | 2023-02-10 | 152.55 | -0.012940796 | 30% |
| A | 2023-02-17 | 148.26 | -0.028121927 | 10% |

3. Display a table showing the top ticker in each decile group

```

df_top_ticker = data_weekly_deciles %>%
  group_by(deciles) %>%
  filter(simple_w_r == max(simple_w_r, na.rm = TRUE)) %>%
  ungroup() %>%
  arrange(desc(deciles)) %>%
  select(tic, simple_w_r, deciles)
df_top_ticker[1:11,]

```

A tibble: 11 × 3

Table 12: Top tickers by decile group

| tic <chr> | simple_w_r <dbl> | deciles <fct> |
|-----------|------------------|---------------|
| SMCI | 0.784176534 | 90% |
| ETN | 0.046986033 | 80% |
| TROW | 0.029170465 | 70% |
| APH | 0.018479909 | 60% |
| PRU | 0.010129310 | 50% |
| AXP | 0.002281286 | 40% |
| NWS | -0.005443235 | 30% |
| IPG | -0.014084507 | 20% |
| WBD | -0.024780176 | 10% |
| ADBE | -0.041523909 | NA |
| NA | NA | NA |

4. Select the top ticker from the 60% decile group

We use this ticker for the rest of the assignment, including in Part 3.

```
summary_table_2_4 = tibble (
  Decile = "60%",
  Ticker = top_ticker_60d
)
summary_table_2_4

top_ticker_60d = as.character(df_top_ticker %>%
  filter(deciles == "60%") %>%
  select(tic)
)
cat("4. The ticker with the highest mean weekly return at the 60% decile is",
top_ticker_60d)
```

A tibble: 1 × 2

| Decile <chr> | Ticker <chr> |
|--------------|--------------|
| 60% | APH |

4. The ticker with the highest mean weekly return at the 60% decile is APH

5. Plot the autocorrelation function for this ticker's entire set of weekly returns

```
# Filter and remove NA values
cscsco_data = data_weekly %>%
  filter(tic == top_ticker_60d) %>%
  na.omit()

acf(cscsco_data$simple_w_r, main = "Autocorrelation Function")
```

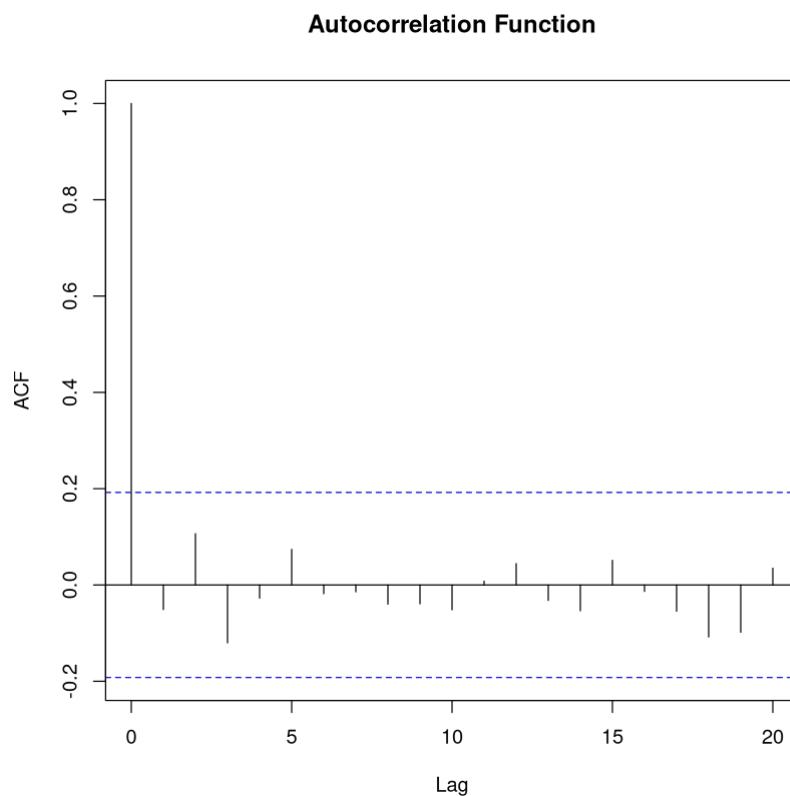


Figure 2: Autocorrelation Function

Part 3 Regression - Fama-French 3 Factor Model

1. Load and clean the weekly Fama-French 3 factor data

```
ff <- read.csv("fama_french_weekly.csv", skip = 4) %>%
  rename(x = X,
    mktrf = Mkt.RF,
    smb   = SMB,
    hml   = HML,
```

```

    rf      = RF) %>%
mutate(
  date   = ymd(as.character(x)),
  mktrf = mktrf / 100,
  smb   = smb / 100,
  hml   = hml / 100,
  rf     = rf / 100
) %>%
transmute(date, mktrf, smb, hml, rf) %>%
filter(!is.na(date)) %>%
arrange(date)

head(ff)

```

A data.frame: 6 × 5

Table 14: Fama-French 3 Factor Model

| | date <date> | mktrf <dbl> | smb <dbl> | hml <dbl> | rf <dbl> |
|---|-------------|-------------|-----------|-----------|----------|
| 1 | 1926-07-02 | 0.0158 | -0.0062 | -0.0086 | 6e-04 |
| 2 | 1926-07-10 | 0.0037 | -0.0090 | 0.0031 | 6e-04 |
| 3 | 1926-07-17 | 0.0098 | 0.0059 | -0.0144 | 6e-04 |
| 4 | 1926-07-24 | -0.0203 | 0.0002 | -0.0017 | 6e-04 |
| 5 | 1926-07-31 | 0.0306 | -0.0189 | -0.0085 | 6e-04 |
| 6 | 1926-08-07 | 0.0204 | 0.0016 | 0.0055 | 6e-04 |

2. Fit the Fama-French 3 factor model to the weekly returns of the stock in Part 2

```

# 1) Get the chosen stock's weekly returns
ticker_data <- data_weekly %>%
  filter(tic == top_ticker_60d) %>%
  select(datadate, simple_w_r) %>%
  filter(!is.na(simple_w_r))

# 2) Join with Fama-French factors (align on week end)
ff_weekly <- ff %>% rename(datadate = date)

merged <- ticker_data %>%
  inner_join(ff_weekly, by = "datadate") %>%
  mutate(excess_return = simple_w_r - rf)

# 3) Fit FF3: excess_return ~ Mkt.RF + SMB + HML
ff3_model <- lm(excess_return ~ mktrf + smb + hml, data = merged)

# 4) Show results

```

```
summary(ff3_model)
```

Call:

```
lm(formula = excess_return ~ mktrf + smb + hml, data = merged)
```

Residuals:

| Min | 1Q | Median | 3Q | Max |
|----------|----------|---------|---------|---------|
| -0.47998 | -0.01072 | 0.00349 | 0.01654 | 0.07078 |

Coefficients:

| | Estimate | Std. Error | t value | Pr(> t) |
|-------------|-----------|------------|---------|--------------|
| (Intercept) | -0.003136 | 0.005454 | -0.575 | 0.566550 |
| mktrf | 1.233157 | 0.315971 | 3.903 | 0.000175 *** |
| smb | 0.368449 | 0.357088 | 1.032 | 0.304698 |
| hml | 0.260722 | 0.320227 | 0.814 | 0.417517 |

Signif. codes: 0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 0.05344 on 98 degrees of freedom

Multiple R-squared: 0.1987, Adjusted R-squared: 0.1742

F-statistic: 8.102 on 3 and 98 DF, p-value: 7.11e-05

Analysis Summary

APH recorded a moderate positive weekly return of around 1.85%, placing it within the 60th performance decile. The autocorrelation analysis revealed no meaningful serial dependence, suggesting that APH's returns behave largely randomly and that past movements offer little predictive power for future performance.

The Fama–French three-factor regression indicated a strong positive sensitivity to overall market returns, with a statistically significant market beta of about 1.23 ($p < 0.001$). The coefficients on size (SMB) and value (HML) factors were positive but not statistically significant, implying a limited relationship with these style factors. The adjusted R^2 of roughly 0.17 suggests that common market factors explain only part of APH's return variation.

Overall, this pattern points to APH's performance being moderately linked to general market movements, while firm-specific drivers continue to play an important role in explaining its excess returns.